



Information Security Policy

OverSiteAI, LLC

Document Version:	2.0
Effective Date:	January 1, 2025
Last Updated:	June 25, 2025
Last Reviewed:	June 27, 2025
Classification:	Restricted
Owner:	Chief Technology Officer
Approved By:	Chief Executive Officer



Table of Contents

1. Purpose and Scope
2. Information Security Principles
3. Roles and Responsibilities
4. Asset Management
5. Access Control
6. Cryptography and Data Protection
7. Physical and Environmental Security
8. Operations Security
9. Communications Security
10. System Development and Maintenance
11. Supplier and Third-Party Security
12. Incident Management
13. Business Continuity and Resilience
14. Compliance Management
15. Security Training and Awareness
16. Policy Management and Governance
17. Enforcement and Accountability
18. Related Policies and Standards
19. Definitions and Glossary
20. Document Control
21. Appendices

Information Security Policy

OversiteAI, LLC

Document Version: 2.0

Effective Date: January 1, 2025

Last Updated: June 25, 2025

Last Reviewed: June 27, 2025



Classification: Restricted

Owner: Chief Technology Officer

Approved By: Chief Executive Officer

1. Purpose and Scope

1.1 Purpose

OversiteAI has established this Information Security Policy as the cornerstone of our commitment to protecting the information assets that drive our business and maintain the trust of our customers. As a software company specializing in client-hosted data collection and correlation solutions, we recognize that security is not just a technical requirement but a fundamental business imperative that shapes how we design, develop, and deliver our products.

Our unique architectural approach, where customer data remains exclusively within customer-controlled environments, represents a deliberate security design choice that eliminates entire categories of risk. This policy establishes the framework for protecting our intellectual property, development infrastructure, and business operations while ensuring that our security practices scale appropriately with our growth from a focused team of under 20 professionals to our future expanded operations.

1.2 Scope

This policy encompasses all aspects of information security at OversiteAI, creating a comprehensive security framework that addresses our fully remote operational model. The policy applies universally to all employees, contractors, consultants, and temporary workers who access our systems or handle our information, regardless of their location or role within the organization.

Our security scope extends beyond traditional boundaries to include all information systems, applications, and infrastructure that we own or operate, recognizing that in a cloud-first architecture, the perimeter is wherever our people and systems operate. This includes our Azure-based development infrastructure, source code repositories, business applications, and the secure home office environments from which our team members work. The policy governs all processes and procedures related to the creation, processing, storage, transmission, and disposal of company information throughout its lifecycle.



1.3 Policy Objectives

Through this policy, OversightAI aims to establish a security program that balances protection with practicality, ensuring that security enhances rather than hinders our ability to innovate and serve our customers. Our primary objective is to protect the confidentiality, integrity, and availability of our information assets, with particular emphasis on our source code and intellectual property that represent the core value of our business.

We are committed to meeting and exceeding all legal, regulatory, and contractual security requirements, viewing compliance not as a burden but as a baseline for building trust with our customers and partners. Our security practices are designed to maintain customer confidence by demonstrating that we apply the same rigorous security standards to our own operations that we enable for our customers through our software solutions.

Beyond protection, this policy aims to minimize business disruption from security incidents through proactive risk management and rapid response capabilities scaled appropriately to our size. We recognize that perfect security is neither achievable nor economically practical for a company of our size, so we focus on risk-based decision making that prioritizes our efforts where they provide the greatest benefit.

Most importantly, this policy seeks to foster a security-conscious culture throughout our organization where every team member understands their role in protecting our assets and feels empowered to raise security concerns without fear of blame or retribution. Security is not the sole responsibility of technical staff but a shared commitment that strengthens with each person's contribution.

2. Information Security Principles

2.1 Core Security Principles

Data Sovereignty and Isolation

OversightAI's fundamental security architecture is built on the principle of complete data isolation between our systems and customer environments. This architectural decision represents our strongest security control, eliminating entire categories of risk that plague traditional SaaS providers. Customer data remains exclusively within customer-controlled infrastructure, processed by our software running on their systems, with no network paths or APIs that would allow data transmission back to our infrastructure.



This isolation is not merely a configuration choice but is enforced through our software architecture, which operates as a fully self-contained application within customer premises. Our development and testing environments maintain this same isolation principle, using only synthetic data that we generate specifically for testing purposes. This approach ensures that even in development, we never handle actual customer data, eliminating risks of accidental exposure or unauthorized access.

We have made the deliberate choice not to use fourth parties or subcontractors for any function that could potentially access customer systems or data. All development, support, and maintenance activities are performed exclusively by direct OversightAI employees who are bound by our security policies and employment agreements. Furthermore, we have eliminated the use of physical media for any data transfer, recognizing that USB drives and similar devices represent an unacceptable risk in our security model.

Defense in Depth

While our architectural isolation provides strong protection, we implement multiple layers of security controls to protect our own infrastructure and intellectual property. This defense in depth strategy ensures that no single control failure can compromise our security posture. We layer technical controls such as encryption and access management with administrative controls like security policies and training, supplemented by physical security measures appropriate to our remote work model.

Each layer of defense is designed to complement the others, creating overlapping protections that increase the effort and skill required for any potential attacker. For example, even if an attacker were to obtain valid credentials, they would still face multi-factor authentication challenges, network access restrictions, and activity monitoring that would detect and respond to suspicious behavior. This layered approach is particularly important for protecting our source code, which represents our primary intellectual property and competitive advantage.

Least Privilege

Access to OversightAI systems and information is granted based on the principle of least privilege, where individuals receive only the minimum access necessary to perform their job functions effectively. This principle applies not just to human users but also to system accounts, applications, and automated processes. We recognize that excessive permissions are not just a security risk but also increase the potential impact of honest mistakes.



Regular access reviews ensure that permissions remain appropriate as roles evolve and projects change. We pay particular attention to privileged access, implementing additional controls and monitoring for administrative accounts that could have broader impact if compromised. Our quarterly review cycle for privileged accounts strikes a balance between security diligence and administrative overhead appropriate to our size.

Continuous Improvement

Security is not a destination but a journey of continuous improvement as threats evolve and our business grows. We maintain a pragmatic approach to security enhancement, focusing on incremental improvements that provide real risk reduction rather than pursuing perfection. Our continuous improvement process is driven by multiple inputs including security assessment results, incident lessons learned, industry threat intelligence, and feedback from our team members who often identify practical security enhancements during their daily work.

We prioritize improvements based on risk reduction potential and implementation effort, ensuring that our limited resources are applied where they provide the greatest benefit. This might mean choosing to implement a simple but effective control immediately rather than waiting for a perfect but complex solution. We document planned improvements in our risk register, allowing us to track progress and demonstrate our commitment to security maturation over time.

2.2 Security by Design

Security considerations are integrated into every phase of our software development lifecycle, from initial concept through deployment and maintenance. This "security by design" approach is more effective and economical than attempting to add security after the fact. During the design phase of any new feature or system, we conduct threat modeling exercises to identify potential security risks and design appropriate controls from the outset.

Our development teams are trained to think like attackers, considering how features might be misused or exploited. Security requirements are defined alongside functional requirements, ensuring that security is not an afterthought but a core component of what we build. We maintain secure coding standards that our developers follow, supplemented by automated security scanning tools integrated into our continuous integration pipeline.



Before any code reaches production, it undergoes security testing appropriate to the change's risk level. This might range from automated vulnerability scanning for minor updates to comprehensive penetration testing for major new features. By building security into our development process, we create products that are secure by default, reducing the burden on our customers to configure security after deployment.

3. Roles and Responsibilities

3.1 Executive Oversight and Strategic Security Governance

Executive Management Team

OversiteAI's Executive Management Team—comprising key corporate officers such as the Chief Executive Officer (CEO), Chief Technology Officer (CTO), and Chief Operating Officer (COO)—bears collective responsibility for the organization's security posture. Executive oversight ensures that information security is integrated into strategic decision-making and aligned with the company's risk tolerance, legal obligations, and business objectives.

The Executive Management Team:

- Endorses and periodically reviews the company's information security policies.
- Allocates sufficient resources, financial and personnel, to support the information security program.
- Receives regular briefings on security matters from the CTO and/or designated security leadership.
- Reviews findings from internal audits, assessments, and incidents with input from security and risk management personnel.

Governance Committees

OversiteAI's security and risk strategy is supported by formal governance bodies such as the **Information Security Steering Committee** and the **Audit & Risk Committee** of the Board of Directors. These entities ensure independent oversight and maintain responsibility for approving critical security decisions and evaluating organizational risk exposure.

Security responsibilities may be delegated to appropriate personnel or committees through documented governance structures.

3.2 Technical Leadership and Implementation

Chief Technology Officer (CTO)



The CTO is accountable for leading OversightAI's information security program as part of broader technology oversight. Operational execution of security strategies may be delegated to specialized personnel or committees while retaining executive visibility and governance.

Key responsibilities include:

- Developing and maintaining security policies, standards, and guidelines in line with business needs and regulatory requirements.
- Ensuring that security is embedded into technology decisions, architecture, and operational planning.
- Leading the response to significant security incidents, with coordination across relevant internal and external stakeholders.
- Providing periodic security metrics and risk updates to the Executive Management Team and Board oversight committees.

Development Team Leadership

Designated engineering leadership, typically the Development Team Lead or an equivalent role, ensures the secure development of OversightAI's software and services.

Key responsibilities include:

- Translating security policies into practical development practices.
- Maintaining secure coding standards and integrating security requirements into software development lifecycles.
- Coordinating application-level security testing and vulnerability remediation.
- Ensuring that development environments are logically separated from production and remain appropriately secured.

DevOps Leadership

The DevOps Lead, or infrastructure lead responsible for platform operations, oversees the secure configuration and operation of OversightAI's infrastructure.

Key responsibilities include:

- Maintaining secure configurations and hardening practices for cloud-based systems (e.g., Azure).
- Managing identity and access management in accordance with least privilege and just-in-time access principles.
- Monitoring infrastructure for threats and misconfigurations using automated tools and responding to alerts.
- Leveraging Infrastructure-as-Code (IaC) and other automation to enforce consistent, auditable security controls.



3.3 Organizational Security Responsibilities

All Employees

Every team member is responsible for supporting OversightAI's security posture by complying with established security policies and participating in the organization's risk-aware culture.

Employee responsibilities include:

- Completing required security training during onboarding and annually thereafter.
- Reporting security incidents or concerns promptly, supported by a no-blame reporting culture.
- Practicing secure behavior such as using strong passwords, locking devices, and verifying suspicious communications.

Security awareness and behavior are monitored and reinforced by periodic simulations, role-based training, and compliance tracking.

Human Resources (HR) Function

The HR function supports OversightAI's security program through secure workforce management practices throughout the employee lifecycle.

Responsibilities include:

- Ensuring appropriate background checks during hiring, proportionate to access and role sensitivity.
- Coordinating onboarding security training and policy acknowledgment prior to system access.
- Managing secure offboarding in coordination with IT, including access revocation and asset return checklists.
- Reinforcing confidentiality and acceptable use expectations through policy communications and annual re-acknowledgments.



4. Asset Management

4.1 Comprehensive Asset Inventory

Effective security begins with knowing what we need to protect. OversightAI maintains a comprehensive inventory of all information assets that could impact our security posture if compromised. This inventory serves multiple purposes beyond security, including license management, budget planning, and disaster recovery preparation. Given our size, we maintain this inventory using Azure's built-in asset management capabilities supplemented by a simple spreadsheet for assets outside Azure's scope.

Our hardware inventory includes all company-provided laptops and mobile devices, tracked with details including serial numbers, assigned users, and security configuration status such as encryption and endpoint protection. For our fully remote workforce, this also includes documentation of home office setups where company data is accessed, though we don't attempt to inventory personal equipment that doesn't touch company data. Each device is tagged with a property identifier that links it to our records, simplifying processes like support requests or incident investigation.

Software asset tracking encompasses both commercial applications and our own developed software. We maintain records of all software licenses to ensure compliance and identify when security updates are available. Our source code repositories receive special attention in asset tracking, as they represent our primary intellectual property. Cloud infrastructure resources are continuously inventoried through Azure's management tools, providing real-time visibility into our infrastructure footprint and ensuring we don't lose track of resources that could become security vulnerabilities if neglected.

4.2 Information Classification Framework

Not all information requires the same level of protection, and attempting to protect everything equally would be both impractical and economically infeasible. Our information classification framework provides clear categories that help employees understand how to handle different types of information appropriately. We deliberately keep our classification scheme simple with just four categories, avoiding the complexity that would make consistent application difficult for a small team.

Confidential information represents our most sensitive data, where unauthorized disclosure could cause serious harm to OversightAI or our stakeholders. This category includes our source code, which embodies our intellectual property and competitive advantage. Security configurations and credentials fall into this category, as their exposure could directly lead to system compromise. Employee personal information,



including salary data and performance reviews, is classified as Confidential to maintain trust and meet privacy obligations. Customer contracts and pricing information also receive this highest classification to preserve business relationships and competitive position.

Restricted information requires protection but with slightly less stringent controls than Confidential data. This includes customer configuration files that, while not containing customer data, could reveal information about customer operations. Internal procedures and runbooks are Restricted to prevent potential attackers from understanding our operations in detail. Project plans and roadmaps fall into this category to protect our competitive strategy while still allowing necessary internal collaboration.

Internal information encompasses general business information that shouldn't be public but doesn't require special protection within the organization. This includes most internal communications, meeting minutes, and general documentation. While we don't want this information published externally, we can share it freely within OversightAI to facilitate collaboration and transparency. The key distinction is that unauthorized disclosure would be embarrassing or inconvenient rather than harmful.

Public information is intended for external consumption and requires no special protection for confidentiality, though we still ensure its integrity to maintain accurate public communications. This includes our marketing materials, public documentation, and information published on our website. Even Public information goes through appropriate review before release to ensure accuracy and consistency with our messaging.

4.3 Asset Lifecycle Management

Every asset follows a lifecycle from acquisition through disposal, and security considerations apply at each stage. When acquiring new assets, whether hardware, software, or cloud services, we evaluate security implications as part of the selection process. This includes reviewing security features, vendor security practices, and how the asset will integrate with our existing security controls. For significant acquisitions, we may require security assessments or contractual security commitments from vendors.

During active use, assets must be properly configured and maintained according to their classification and security requirements. Hardware devices receive standard security configurations before deployment, including encryption, endpoint protection, and remote management capabilities. Software assets are kept current with security



patches according to our patch management procedures. Cloud resources are configured according to our security baselines, with infrastructure-as-code ensuring consistent and auditable configurations.

Asset disposal requires particular attention to ensure that sensitive information doesn't leave our control. For hardware devices, we use encryption to ensure that data remains protected even if devices are lost or improperly disposed. When devices reach end-of-life, we perform secure wiping before disposal or recycling. For cloud resources, we carefully remove all data and configurations before releasing resources. Software licenses are properly transferred or terminated to maintain compliance and prevent unauthorized use.

4.4 Acceptable Use Standards

Clear acceptable use standards help employees understand their responsibilities when using company assets. We maintain a balance between enabling productivity and protecting security, avoiding overly restrictive policies that employees might circumvent. Our acceptable use policy emphasizes personal responsibility and good judgment rather than attempting to enumerate every possible scenario.

Company assets should be used primarily for business purposes, though we recognize that some personal use is inevitable and even beneficial for work-life balance in a remote environment. We draw the line at uses that could compromise security, violate laws, or damage our reputation. This includes prohibitions on installing unauthorized software that could introduce vulnerabilities, using company systems for personal business ventures, or accessing inappropriate content that could introduce malware or create hostile work environment claims.

We particularly emphasize the risks of mixed personal and business use on devices. While we don't prohibit reasonable personal use like checking personal email or reading news during breaks, employees must understand that company devices are subject to monitoring and investigation if security incidents occur. We strongly discourage storing personal files on company devices to avoid awkward situations during investigations or device returns. For critical security tools like password managers, we provide company licenses to avoid the temptation to use personal accounts that might not meet our security standards.



5. Access Control

5.1 Access Management Philosophy

Access control forms the foundation of our security program, determining who can access what resources under which circumstances. Our approach to access management balances security requirements with operational efficiency, recognizing that overly complex access controls can lead to workarounds that undermine security. We implement role-based access control (RBAC) wherever possible, simplifying administration and ensuring consistency while maintaining appropriate segregation of duties for critical functions.

Our access control philosophy is grounded in the principle of least privilege, but we implement this pragmatically rather than dogmatically. We start by defining broad role-based access that covers most common scenarios, then layer additional restrictions for particularly sensitive resources. This approach avoids the administrative burden of managing highly granular permissions while still protecting our most critical assets. Regular access reviews ensure that permissions remain appropriate as roles evolve and projects change.

We emphasize preventive controls supported by detective measures, recognizing that in a small organization, we cannot maintain real-time monitoring of all access. Instead, we focus on strong authentication, clear authorization models, and comprehensive logging that enables investigation when concerns arise. This balanced approach provides effective security without requiring a dedicated security operations center that would be impractical at our size.

5.2 Identity and Authentication Management

User Identity Lifecycle

Every user's digital identity at OversightAI follows a carefully managed lifecycle from creation through eventual deactivation. New accounts are created only through a formal request and approval process that documents the business justification and required access levels. We use standardized role templates that define common permission sets, reducing errors and ensuring consistency. Each user receives a unique identifier that remains constant throughout their tenure, enabling accurate audit trails even as names or roles change.

During onboarding, new users receive their initial credentials through a secure channel and must change default passwords on first login. We enforce technical controls that prevent password reuse and require passwords meeting our complexity standards. The



onboarding process includes security training that helps users understand not just the technical requirements but the reasoning behind them, fostering a security-conscious mindset from day one.

Account maintenance throughout the employment lifecycle includes regular reviews to ensure access remains appropriate. When employees change roles, we follow a revoke-and-re-grant process rather than simply adding new permissions, preventing the accumulation of excessive access over time. Temporary elevated access for specific projects is time-bound and automatically expires, requiring re-justification for extension. This lifecycle management ensures that access privileges accurately reflect current responsibilities rather than historical accumulations.

Strong Authentication Implementation

Password complexity alone is insufficient protection in today's threat landscape, so we implement comprehensive authentication controls that balance security with usability. Our password policy requires a minimum of 14 characters combining multiple character types, but we emphasize passphrase approaches that are both stronger and easier to remember than traditional complex passwords. We prohibit passwords containing dictionary words or personal information that could be discovered through social engineering or public records research.

Multi-factor authentication (MFA) is mandatory for all access to systems containing Confidential or Restricted information. We've standardized on Azure's MFA capabilities, which integrate seamlessly with our infrastructure while providing flexibility in authentication methods. Users can choose between mobile app notifications, SMS codes (though we discourage this due to SIM swapping risks), or hardware tokens for those requiring the highest security. This flexibility ensures that MFA enhances rather than hinders productivity.

Account lockout policies protect against brute force attacks while avoiding denial-of-service scenarios where attackers could intentionally lock out legitimate users. After five failed login attempts, accounts enter a progressive lockout state with increasing delays between allowed attempts. This slows attackers while allowing legitimate users who simply forgot their passwords to eventually succeed. Security alerts notify administrators of lockout events, enabling investigation of potential attacks while supporting users who need assistance.

5.3 Authorization and Privileged Access

Role-Based Authorization Model



Our authorization model maps permissions to job functions rather than individuals, simplifying administration and ensuring consistency. We've defined standard roles that cover most access needs such as Developer, DevOps Engineer, Business Analyst, and Administrative Staff. Each role includes predefined access to systems and data necessary for that function. This role-based approach means new employees can be productive quickly while maintaining appropriate security boundaries.

For scenarios requiring access beyond standard roles, we implement an exception process that documents the business justification and includes appropriate approval. These exceptions are time-limited by default, requiring periodic review and re-approval to continue. This prevents the accumulation of special permissions that often creates security vulnerabilities in growing organizations. All exceptions are tracked centrally, enabling quick identification of users with elevated access during security reviews or incident investigations.

Segregation of duties is implemented for critical functions where no single person should have complete control. For example, code development and production deployment require different roles, ensuring that changes receive appropriate review before reaching production. Financial functions separate authorization and execution roles, preventing unauthorized transactions. While our small size requires some individuals to wear multiple hats, we ensure that conflicting roles are not combined in ways that would eliminate important controls.

Privileged Access Management

Administrative access receives special attention given its potential for widespread impact if misused or compromised. We maintain separate administrative accounts distinct from regular user accounts, ensuring that administrative privileges are used only when necessary. These privileged accounts follow naming conventions that clearly identify them as administrative, reducing the risk of accidental use for routine activities that don't require elevated permissions.

Just-in-time (JIT) access principles guide our privileged access management, though we implement this pragmatically given our size. Rather than maintaining standing administrative access, we grant time-bound elevated permissions when needed for specific tasks. Azure Privileged Identity Management (PIM) automates much of this process, allowing users to request and activate privileged roles for defined periods. This approach reduces our attack surface by ensuring privileged access exists only when actively needed.



Monitoring and auditing of privileged access provides detective controls to complement our preventive measures. All privileged account activities are logged with enhanced detail, including commands executed and data accessed. These logs feed into our security monitoring systems with alerts for unusual patterns such as privileged access outside business hours or from unexpected locations. Quarterly reviews of privileged access ensure that administrative permissions remain appropriate and that any standing privileged access is truly justified by operational requirements.

5.4 Remote Access Security

As a fully remote organization, secure remote access is not an exception but our standard operating model. All access to OversightAI development resources requires connection through our VPN infrastructure, which provides encrypted tunnels and enables consistent security policy enforcement regardless of user location. The VPN serves not just as an encryption mechanism but as a policy enforcement point where we verify device compliance before allowing connections.

Device health checks confirm that connecting systems meet our security baselines including current patches, enabled encryption, and active endpoint protection. Non-compliant devices are redirected to a remediation portal that guides users through required updates before access is granted. This automated enforcement ensures consistent security without requiring manual intervention for routine cases. We maintain exemption processes for exceptional situations, but these require security team review and are time-limited.

Session management controls limit the impact of compromised sessions or unattended devices. VPN sessions timeout after 30 minutes of inactivity, requiring re-authentication to continue. This balances security with usability, protecting against forgotten sessions while not interrupting active work. All remote access is logged with details including connection times, source locations, and resources accessed. These logs support both security investigations and compliance requirements while respecting employee privacy by focusing on access patterns rather than specific activities.



6. Cryptography and Data Protection

6.1 Encryption Strategy

Cryptographic controls provide our last line of defense, ensuring that even if other controls fail, sensitive data remains protected. Our encryption strategy addresses data throughout its lifecycle, from creation through transmission, storage, and eventual destruction. We standardize on industry-proven algorithms and implementations rather than attempting to create custom cryptographic solutions, recognizing that cryptography is a specialized field where amateur attempts often introduce vulnerabilities.

For data at rest, we implement AES-256 encryption universally for all systems storing Confidential information. This includes full disk encryption on all endpoints using BitLocker for Windows systems and FileVault for macOS devices. Azure Storage Service Encryption protects our cloud storage automatically, providing transparent encryption without performance impact. Database encryption uses Transparent Data Encryption (TDE) where available, ensuring that database files and backups remain protected even if somehow exfiltrated.

Data in transit receives equal protection through mandatory TLS 1.2 or higher for all network communications. We've disabled older protocols that have known vulnerabilities, even though this occasionally causes compatibility challenges with legacy systems. Our VPN infrastructure adds an additional encryption layer for remote access, creating defense in depth for our most sensitive communications. Email encryption is available for sensitive external communications, though we prefer secure file sharing platforms that provide better control and audit capabilities.

6.2 Key Management Practices

Effective key management is essential for maintaining the value of encryption, as improperly managed keys can render even strong encryption useless. We leverage Azure Key Vault as our central key management system, providing hardware security module (HSM) protection for our most sensitive keys while maintaining operational efficiency. This cloud-native approach eliminates the complexity of managing physical HSMs while providing comparable security for our needs.

Key rotation follows defined schedules based on key usage and sensitivity. Master keys rotate annually, while data encryption keys rotate more frequently based on usage volume. Automated rotation processes ensure consistency and prevent human error, with Azure Key Vault handling much of the complexity transparently. We maintain careful documentation of rotation schedules and procedures, ensuring continuity even



if key personnel are unavailable during critical rotation windows.

Split knowledge and dual control principles protect our most critical keys, ensuring no single person can compromise core encryption systems. Master key operations require two authorized individuals working together, with neither having complete access independently. While this adds operational overhead, we've limited its application to truly critical keys where the additional protection justifies the complexity. For routine encryption operations, we balance security with operational efficiency through role-based access controls and comprehensive audit logging.

6.3 Certificate Management

Digital certificates enable secure communications and authentication throughout our infrastructure, but they require careful management to maintain their security value. We obtain certificates exclusively from trusted Certificate Authorities (CAs), avoiding self-signed certificates except in isolated development environments where they pose no production risk. Our standardized certificate request process ensures consistency and includes security review for certificates protecting critical services.

Certificate lifecycle management prevents the service disruptions and security vulnerabilities that can result from expired certificates. We maintain a central inventory of all certificates with expiration tracking and automated alerts at 90, 60, and 30 days before expiration. This graduated alerting ensures adequate time for renewal while escalating urgency as expiration approaches. For critical services, we maintain documented renewal procedures that can be executed by multiple team members, preventing single points of failure.

Private key protection receives special attention given that certificate security depends entirely on private key confidentiality. Keys are generated using appropriate randomness and immediately protected with access controls limiting availability to necessary systems and personnel. We never transmit private keys via email or other insecure channels, instead using secure transfer mechanisms when key movement is absolutely necessary. When certificates are retired or replaced, we ensure proper revocation and secure destruction of associated private keys.



7. Physical and Environmental Security

7.1 Remote Work Security Model

As a fully distributed organization, OversightAI's approach to physical security differs significantly from traditional office-based companies. Rather than securing a central facility, we focus on helping employees establish and maintain secure home office environments. This distributed security model recognizes that each employee's home becomes an extension of our security perimeter, requiring thoughtful controls that respect both security needs and personal privacy.

Our remote work security standards establish baseline requirements while acknowledging the diversity of home environments. Employees must maintain dedicated workspace areas where business activities can be conducted without unauthorized observation or interference. This doesn't require a separate room but does mean positioning screens away from windows or common areas where sensitive information might be observed. We provide privacy screens for laptops to enable secure work in shared spaces when necessary, recognizing that absolute isolation isn't always practical in home environments.

Physical access controls in home offices focus on preventing unauthorized access to company devices and information. Employees must secure devices when not in use, either through physical locks or secure storage. We provide cable locks for desktop systems and expect laptops to be stored securely when not actively supervised. While we don't mandate safes or specialized storage, we do require that devices never be left unattended in vehicles or public spaces where theft risk is elevated. These practical measures significantly reduce risk without imposing unreasonable burdens on employees' home lives.

7.2 Device Security Controls

Every device that accesses OversightAI information becomes a potential entry point for attackers, making device security a critical control. Our device security program begins before deployment, with standard security configurations applied to all company-provided equipment. These configurations include full disk encryption, endpoint detection and response software, automatic screen locking, and remote wipe capabilities. We maintain configuration templates that ensure consistency while allowing flexibility for different roles' requirements.

Travel security receives particular attention given that mobile devices face elevated risks outside the relatively controlled home office environment. Employees must carry devices as carry-on luggage when flying, preventing both theft and potential tampering



in checked baggage. Devices should never be left unattended in public spaces, even briefly, as theft can occur in seconds. When working in public spaces, employees must use privacy screens and position themselves to prevent shoulder surfing of sensitive information. International travel requires additional precautions including travel-specific devices when visiting high-risk countries.

Lost or stolen device procedures emphasize rapid response to minimize potential impact. Employees must report device loss within two hours of discovery, enabling rapid remote wipe before potential data access. We maintain 24/7 contact procedures for device loss reports, recognizing that incidents don't follow business hours. Post-incident reviews examine circumstances to identify process improvements, but we maintain a no-blame culture for good-faith losses, recognizing that punitive approaches discourage prompt reporting that enables effective response.

7.3 Clean Desk and Information Handling

The clean desk policy extends beyond traditional desk cleanliness to encompass comprehensive information handling practices in home offices. When not actively in use, all confidential information must be secured from unauthorized access. This includes both digital and physical forms, recognizing that printouts and handwritten notes can expose sensitive information just as effectively as unsecured screens. Employees should minimize printing of sensitive information, but when necessary, must store printouts in locked drawers or filing cabinets when not actively referenced.

Screen locking is mandatory whenever leaving the workspace, even briefly. We've configured automatic locking after 10 minutes of inactivity as a backup, but employees should manually lock screens whenever stepping away. This habit is particularly important in home environments where family members or visitors might inadvertently see sensitive information. Password-protected screensavers provide an additional layer of protection while maintaining quick access for returning users.

Information disposal requires equal attention to prevent dumpster diving or casual discovery of sensitive information. All confidential documents must be shredded before disposal, with cross-cut shredders providing adequate protection for most materials. For highly sensitive documents, we maintain a service for certified destruction with chain of custody documentation. Electronic media disposal follows secure wiping procedures detailed in our asset management processes. Even seemingly innocuous items like sticky notes or whiteboard content must be properly cleared if they contained sensitive information.



8. Operations Security

8.1 Architectural Security Through Customer Data Isolation

The cornerstone of OversightAI's operational security is our fundamental architectural decision to maintain complete separation between customer data and our systems. This isn't merely a configuration choice or policy decision but is built into the very architecture of our solutions. Our software operates as a fully self-contained application deployed entirely within customer infrastructure, with no phone-home capabilities, telemetry collection, or data synchronization back to our systems.

This architectural isolation eliminates entire categories of operational security concerns that plague traditional SaaS providers. We cannot suffer a breach that exposes multiple customers' data because we simply don't have access to that data. Customer data is processed exclusively within infrastructure owned and controlled by each customer, with our software serving as the processing engine but never the data custodian. This design philosophy extends throughout our operational practices, influencing everything from how we provide support to how we develop and test new features.

Our development and testing environments mirror this isolation, using only synthetic data specifically generated for testing purposes. This synthetic data is carefully crafted to exercise all code paths and edge cases without ever incorporating real customer information. Even when customers report issues, our troubleshooting occurs through log analysis and configuration reviews rather than data access. This approach occasionally makes debugging more challenging, but we view this as a worthwhile tradeoff for the security benefits it provides both to our customers and to OversightAI by reducing our regulatory exposure and breach risk.

8.2 Change Management and Configuration Control

Disciplined change management prevents many security incidents by ensuring that modifications to our systems are properly planned, reviewed, and tested before implementation. Our change management process scales appropriately to our size, avoiding bureaucratic overhead while maintaining necessary controls. All changes to production systems follow a defined workflow that includes impact assessment, approval requirements, testing validation, and rollback planning.

Change proposals begin with clear documentation of what will be modified and why, including specific configuration changes or code deployments planned. Security impact assessment is integrated into this proposal process, with explicit consideration of how changes might affect our security posture. This might range from simple



acknowledgment that security impact was considered and found minimal, to detailed analysis of new attack surfaces or control modifications for security-significant changes.

Testing in non-production environments is mandatory before production deployment, with test environments maintaining reasonable parity to production to ensure valid results. Our infrastructure-as-code approach enables rapid provisioning of test environments that accurately reflect production configurations. After testing validates both functionality and security, changes deploy through automated pipelines that ensure consistency and maintain audit trails. Post-implementation reviews for significant changes capture lessons learned and identify process improvements.

8.3 Security Monitoring and Event Management

Effective security monitoring enables detection of potential incidents before they escalate into breaches. We leverage Azure's native security monitoring capabilities supplemented by additional tools where necessary, providing comprehensive visibility into our environment without the complexity of managing separate SIEM infrastructure. Azure Security Center serves as our primary security monitoring dashboard, aggregating alerts from across our infrastructure and providing actionable recommendations for improvement.

Log collection encompasses security-relevant events from all systems, including authentication attempts, privilege usage, configuration changes, and data access. We've configured logging levels to capture necessary detail for investigation while avoiding log bloat that would make analysis impractical. Centralized log storage in Azure Log Analytics provides both real-time analysis capabilities and long-term retention for compliance and forensic purposes. We maintain logs for one year by default, with extended retention for specific compliance requirements.

Alert tuning balances detection sensitivity with operational practicality. We've learned that too many false positive alerts lead to alert fatigue and missed genuine incidents. Through iterative tuning, we've developed alert rules that reliably detect concerning activities while maintaining manageable alert volumes. Critical alerts such as multiple failed authentication attempts or privilege escalation trigger immediate notification, while lower-priority alerts aggregate into daily summaries for review. This tiered approach ensures rapid response to urgent issues while maintaining visibility into lower-level concerns.



8.4 Vulnerability Management Program

Proactive vulnerability management reduces our attack surface by identifying and remediating weaknesses before they can be exploited. Our vulnerability management program combines multiple detection methods including automated scanning, manual testing, and threat intelligence monitoring. We've sized this program appropriately for our organization, focusing on high-impact activities rather than attempting comprehensive coverage that would require dedicated security staff.

Azure Security Center provides continuous vulnerability assessment for our cloud infrastructure, automatically scanning for misconfigurations and missing patches. These automated scans run continuously, providing near real-time visibility into our vulnerability posture. For application security, we integrate vulnerability scanning into our CI/CD pipeline, catching many issues before they reach production. Static application security testing (SAST) tools analyze our source code for common vulnerability patterns, while software composition analysis (SCA) identifies known vulnerabilities in third-party components we use.

Remediation prioritization follows a risk-based approach that considers both vulnerability severity and actual exploitability in our environment. Critical vulnerabilities in internet-facing systems receive immediate attention, with patches applied within 7 days or compensating controls implemented if patching must be delayed. High and medium severity vulnerabilities follow 30 and 90-day remediation timelines respectively, with exceptions requiring documented risk acceptance. This structured approach ensures consistent treatment while maintaining flexibility for operational necessities.

8.5 Malware Protection Strategy

While our Linux-based infrastructure and macOS development environments face different malware threats than traditional Windows enterprises, protection against malicious software remains important. Our malware protection strategy layers multiple controls, recognizing that no single solution provides complete protection. Endpoint detection and response (EDR) software on all endpoints provides real-time protection against known malware while also detecting suspicious behaviors that might indicate novel threats.

Email serves as a primary malware delivery vector, so we implement comprehensive email security including attachment sandboxing and link protection. Suspicious attachments execute in isolated environments to detect malicious behavior before delivery to users. URL rewriting enables time-of-click protection against links that become malicious after delivery. Combined with user training on recognizing



suspicious emails, these technical controls significantly reduce malware risk from email-based attacks.

Regular security awareness training helps employees recognize and avoid malware delivery attempts. Rather than relying solely on technical controls, we empower employees to be part of our defense. Training covers recognition of phishing attempts, safe handling of unexpected attachments, and proper procedures when malware infection is suspected. We conduct periodic phishing simulations to assess training effectiveness and identify employees who may need additional support. This human-centric approach complements our technical controls, creating defense in depth against malware threats.

9. Communications Security

9.1 Network Security Architecture

Our network security architecture reflects our cloud-first, remote-work reality, where traditional perimeter-based security models no longer apply. Instead, we implement a zero-trust approach that verifies every connection regardless of source. Azure's native network security capabilities provide the foundation, with Network Security Groups (NSGs) enforcing micro-segmentation between different application tiers and functions. This granular control ensures that even if one segment is compromised, lateral movement to other systems is restricted.

Network traffic flows through defined paths with inspection and logging at key points. Azure Firewall provides centralized security policy enforcement and threat intelligence-based filtering. Rather than maintaining complex firewall rules, we leverage Azure's managed security services that automatically update with the latest threat intelligence. This approach provides enterprise-grade network security without requiring dedicated security operations staff to maintain rule sets and monitor for emerging threats.

For our remote workforce, all connections to development resources traverse our VPN infrastructure, which provides not just encryption but also consistent security policy enforcement. The VPN concentrator performs deep packet inspection, identifying and blocking malicious traffic patterns. We've implemented split-tunneling configurations that route only corporate traffic through the VPN, improving performance while maintaining security for sensitive connections. Regular network security assessments validate our configurations and identify any drift from security baselines.



9.2 Application Security Framework

Security is embedded throughout our application development lifecycle rather than bolted on as an afterthought. Our secure development framework begins with threat modeling during design phases, where we systematically identify potential attack vectors and design appropriate mitigations. These threat models are living documents, updated as features evolve and new threats emerge. By considering security from the earliest stages, we avoid the costly rework that comes from discovering vulnerabilities late in development.

Secure coding standards provide consistent guidance to developers, covering common vulnerability patterns and language-specific security considerations. These standards are practical rather than academic, with real examples from our codebase showing both vulnerable and secure implementations. We maintain these standards in our development wiki, making them easily accessible and searchable. Regular secure coding training ensures developers understand not just what the standards require but why these practices matter for security.

Security testing is integrated into our continuous integration pipeline, providing rapid feedback on potential vulnerabilities. Static analysis tools scan every code commit, flagging potential security issues before they can be merged. Dynamic analysis during integration testing exercises application security controls, verifying that authentication, authorization, and input validation work as designed. For major releases, we engage penetration testers who bring fresh perspectives and specialized skills to challenge our assumptions and find issues our internal processes might miss.

9.3 Secure Information Transfer

Information must often move between systems and organizations, creating potential exposure points that require careful control. Our secure information transfer mechanisms provide appropriate protection based on data sensitivity while maintaining usability. For routine internal transfers, our standard encrypted channels provide adequate protection. For sensitive external transfers, we leverage additional controls including dedicated secure file transfer platforms that provide encryption, access control, and audit trails.

Large data transfers receive special scrutiny, as they can indicate either legitimate business needs or potential data exfiltration. Transfers exceeding defined thresholds require approval and documentation of business purpose. We monitor for unusual patterns such as large transfers outside business hours or to unusual destinations. This monitoring is tuned to our normal business patterns, avoiding false positives while detecting genuinely concerning activities. When investigation reveals legitimate



transfers, we update our baselines to prevent future false alerts.

Data loss prevention (DLP) controls provide an additional layer of protection against inadvertent or malicious data exposure. Azure Information Protection enables classification and protection of sensitive documents, with policies that prevent unauthorized sharing. While we avoid overly restrictive DLP policies that impede legitimate work, we do implement targeted controls for our most sensitive information such as source code and security configurations. These controls alert on potential exposures while maintaining audit trails for investigation and compliance purposes.

10. System Development and Maintenance

10.1 Secure Development Lifecycle

Security is not a phase of development but a continuous consideration throughout our software development lifecycle. From initial concept through deployment and maintenance, security requirements receive equal priority with functional requirements. This integrated approach produces more secure software while actually reducing overall development time by catching issues early when they're less expensive to fix.

Requirements gathering explicitly includes security requirements derived from threat modeling and compliance obligations. These requirements are specific and testable rather than vague aspirations. For example, instead of simply requiring "secure authentication," we specify multi-factor authentication support, session timeout parameters, and password complexity requirements. This specificity ensures developers understand expectations and can build appropriate solutions from the start rather than retrofitting security after functional development.

Design reviews include security architecture evaluation, ensuring that security controls are properly integrated rather than superficially applied. We look for defense in depth, ensuring that no single control failure can compromise security. The principle of least privilege guides design decisions, with components receiving only the permissions necessary for their function. By addressing security during design, we avoid architectural flaws that would be difficult or impossible to fix later without major rework.



10.2 Code Security Practices

Writing secure code requires both knowledge and discipline. Our developers receive regular training on secure coding practices specific to the languages and frameworks we use. This training goes beyond academic vulnerability discussions to provide practical guidance on avoiding common pitfalls. We maintain a library of secure coding patterns that developers can reference and reuse, reducing the likelihood of introducing vulnerabilities through inexperience or oversight.

Code reviews include explicit security checks alongside functional correctness validation. Reviewers use security-focused checklists that highlight common vulnerability patterns to watch for. These reviews serve dual purposes: catching security issues before they reach production and providing continuous security education as developers learn from reviewing others' code. We've found that peer review often catches subtle security issues that automated tools miss, particularly logic flaws that require understanding business context.

Third-party components receive particular scrutiny, as they can introduce vulnerabilities outside our direct control. We maintain an inventory of all third-party libraries and frameworks used in our applications, tracking version information and known vulnerabilities. Automated tools alert us when new vulnerabilities are discovered in components we use, enabling rapid assessment and remediation. We prefer widely-used, actively-maintained components over obscure alternatives, as popular projects typically receive more security scrutiny and faster patches.

10.3 Testing and Quality Assurance

Security testing is integrated throughout our testing processes rather than relegated to a final gate before release. Unit tests include security-focused test cases that verify authentication, authorization, and input validation logic. These automated tests run with every build, providing immediate feedback when changes break security controls. Integration tests exercise security controls in realistic scenarios, verifying that components work together securely.

Penetration testing provides valuable external validation of our security controls. We engage qualified penetration testers annually for comprehensive assessments, with additional targeted tests for major new features or architectural changes. These testers bring specialized skills and fresh perspectives, often finding issues that internal testing missed. We treat penetration test findings as learning opportunities rather than failures, using them to improve both our products and our development processes.



Test environments maintain security controls appropriate to the data they handle. While test environments use only synthetic data, we still implement access controls and monitoring to prevent them from being used as launching points for attacks on production systems. Test environment configurations closely mirror production to ensure security testing results are valid. After testing completes, we sanitize test environments to prevent accumulation of test data that might be mistaken for real information.

10.4 Change Control and Release Management

Every change to production systems follows our formal change control process, which scales appropriately to change risk and complexity. Minor patches might follow an expedited process with abbreviated testing, while major releases undergo comprehensive review and testing. All changes require documented approval from appropriate stakeholders, creating accountability and ensuring changes align with business objectives.

Security impact assessment is mandatory for all changes, though the depth varies with change scope. A simple configuration update might require only a brief statement that security impact was considered and found minimal. Major architectural changes require detailed analysis of new attack surfaces, modified trust boundaries, and altered control effectiveness. This graduated approach ensures appropriate scrutiny without creating bureaucratic barriers to routine maintenance.

Release management procedures ensure consistent, repeatable deployments that maintain security throughout the process. Automated deployment pipelines enforce security checks, preventing deployment of code that fails security scans or lacks required approvals. Rollback procedures are tested regularly, ensuring we can quickly revert problematic changes. Post-deployment validation confirms that security controls operate correctly in the production environment, catching any issues that might arise from environmental differences.



11. Supplier and Third-Party Security

11.1 Third-Party and Open Source Risk Management

Every third-party relationship introduces potential security risks that must be understood and managed. Our vendor risk management program scales vendor scrutiny to the risk they present, avoiding wasteful deep-dives on low-risk vendors while ensuring appropriate diligence for critical suppliers. Initial vendor assessment considers factors including the type of data they might access, their role in our operations, and their security maturity.

Security assessment begins during vendor selection, not after contracts are signed. We include security criteria in our evaluation process, considering vendors' security certifications, incident history, and willingness to meet our security requirements. For vendors who will handle sensitive data or provide critical services, we may require completion of security questionnaires or evidence of third-party security audits. This upfront assessment helps avoid the frustration of discovering security gaps after committing to a vendor.

Contractual security requirements translate our security expectations into binding obligations. Standard security clauses address common requirements like data protection, incident notification, and compliance with applicable laws. For higher-risk vendors, we negotiate specific security controls and audit rights. While we're realistic about our negotiating leverage as a small company, we've found that many vendors willingly accommodate reasonable security requirements when asked. Where vendors cannot meet our ideal requirements, we document compensating controls or risk acceptance decisions.

11.2 Cloud Service Provider Management

Microsoft Azure serves as our primary infrastructure provider, making their security posture critical to our own. We've chosen Azure partly for their comprehensive security capabilities and compliance certifications, which provide assurance without requiring detailed audits beyond our capabilities. Regular review of Azure's security documentation and compliance reports ensures their controls remain aligned with our needs.

Beyond trusting Azure's baseline security, we actively configure and monitor the security of our Azure resources. The shared responsibility model means that while Azure secures the underlying infrastructure, we must properly configure and use the services they provide. We leverage Azure Security Center's recommendations to identify and remediate misconfigurations. Regular reviews of our Azure architecture



ensure we're following cloud security best practices and taking advantage of new security features as they become available.

Our cloud exit strategy, while hopefully never needed, ensures we're not locked into Azure in ways that could compromise our security or business flexibility. We maintain infrastructure-as-code definitions that could be adapted to other cloud providers if necessary. Regular exports of critical configurations and data ensure we could recover operations elsewhere if required. This planning exercise also improves our disaster recovery capabilities and helps identify unnecessary dependencies that could be eliminated to improve portability.

11.3 Supply Chain Security

Software supply chain attacks have become increasingly common, making the security of our development tools and dependencies critical. We carefully evaluate the security of development tools, preferring established vendors with strong security track records. Development tool access is restricted to authorized developers, with additional controls for tools that could modify our code or build processes. Regular updates ensure we benefit from security patches while testing in non-production environments first to avoid breaking changes.

Dependency management receives particular attention given the extensive use of open-source components in modern software development. We maintain a complete inventory of all third-party libraries and frameworks used in our products, tracking not just direct dependencies but also transitive dependencies. Automated tools scan for known vulnerabilities and alert us to necessary updates. We prefer widely-used, actively-maintained components over obscure alternatives, as popular projects typically receive more security scrutiny.

Build pipeline security ensures that our software cannot be compromised during the build and deployment process. Build systems are isolated from production environments and developer workstations, reducing the risk of compromise from either direction. All build artifacts are signed, providing assurance that deployed code matches what we built. Build logs provide audit trails of what was built, when, and by which processes. These controls protect against sophisticated attacks that might attempt to inject malicious code during the build process.



12. Incident Management

12.1 Incident Response Framework

Despite our best preventive efforts, security incidents may still occur. Our incident response framework ensures rapid, effective response that minimizes impact while preserving evidence for investigation and improvement. The framework scales to incident severity, avoiding over-reaction to minor issues while ensuring major incidents receive appropriate resources. Clear roles and responsibilities prevent confusion during the stress of incident response.

Incident classification guides response efforts, with categories ranging from minor security events requiring only documentation to critical incidents demanding immediate executive attention. This classification considers multiple factors including data sensitivity, system criticality, and potential for escalation. By establishing classification criteria in advance, we avoid debates during incidents about appropriate response levels. Regular review of classifications ensures they remain aligned with our evolving risk profile.

Our incident response team structure reflects our small size while ensuring necessary capabilities. Rather than maintaining a dedicated incident response team, we've trained key technical staff in incident response procedures and rotate on-call responsibilities. The CTO serves as default incident commander for security incidents, with defined alternates for coverage. This approach provides incident response capability without the overhead of a dedicated team that would be underutilized in a small organization.

12.2 Detection and Reporting

Rapid incident detection enables faster response and reduced impact. Our detection capabilities combine automated monitoring with human observation, recognizing that each has strengths the other lacks. Automated systems excel at detecting known attack patterns and anomalies from baselines, while humans often notice subtle indicators that don't trigger automated alerts. We cultivate a security-aware culture where employees feel empowered to report concerns without fear of being wrong.

Reporting procedures emphasize speed over perfection, encouraging immediate notification even when full details aren't yet known. We maintain multiple reporting channels including email, phone, and chat to ensure availability regardless of circumstances. After-hours procedures ensure incidents outside business hours receive appropriate attention. The on-call engineer performs initial triage, escalating to additional resources as needed. This graduated response balances availability with sustainability for our small team.



Initial response focuses on containment and evidence preservation rather than immediate remediation. We've learned that hasty cleanup efforts often destroy evidence needed for investigation and can even worsen incidents by triggering additional attacker actions. Our response procedures emphasize documenting observations and preserving system state before making changes. This disciplined approach has proven valuable in post-incident analysis and has occasionally revealed that apparent incidents were actually legitimate but unusual activities.

12.3 Investigation and Recovery

Incident investigation follows structured procedures that ensure thoroughness while maintaining chain of custody for potential legal proceedings. We document all actions taken during investigation, including commands run, files accessed, and conclusions drawn. This documentation serves multiple purposes: enabling others to understand and verify our work, providing evidence for any legal proceedings, and creating training materials for future incidents.

Technical investigation leverages the comprehensive logging we maintain across our environment. Centralized log storage enables correlation of events across systems, often revealing attack patterns not visible from individual systems. We maintain investigation playbooks for common incident types, providing step-by-step guidance that ensures consistency and completeness. These playbooks are living documents, updated after each incident with lessons learned and new techniques discovered.

Recovery procedures prioritize business continuity while ensuring security is maintained or improved. We resist the temptation to simply restore systems to their pre-incident state if that state enabled the incident. Instead, recovery includes implementing additional controls or configuration changes to prevent recurrence. Post-recovery validation confirms not just that systems function correctly but that security controls are properly implemented. Extended monitoring after recovery watches for incident recurrence or related attacks.

12.4 Post-Incident Activities

Every incident, regardless of outcome, provides learning opportunities that strengthen our security posture. Post-incident reviews examine not just technical aspects but also process effectiveness and team performance. We maintain a blame-free culture for these reviews, focusing on systemic improvements rather than individual failures. This approach encourages honest discussion and surfaces issues that might otherwise remain hidden.



Lessons learned are documented and drive concrete improvements to our security controls, processes, and training. We track these improvements to closure, ensuring that identified weaknesses are actually addressed rather than simply documented. Metrics from incidents feed into our risk management process, helping prioritize security investments based on actual threat patterns rather than theoretical risks.

Communication about incidents balances transparency with appropriate confidentiality. Internal stakeholders receive detailed briefings appropriate to their roles, ensuring organizational learning while maintaining need-to-know principles. External communications, when required by regulation or contract, provide necessary information while protecting sensitive details that might enable copycat attacks. We maintain template communications for common scenarios, enabling rapid response while ensuring consistency and completeness.

13. Business Continuity and Resilience

13.1 Business Continuity Planning

Business continuity planning ensures OversightAI can maintain essential operations despite disruptive events, whether security incidents, natural disasters, or other interruptions. Our planning process begins with business impact analysis that identifies critical business functions and their supporting resources. For a software company like ours, this primarily focuses on maintaining development capabilities, protecting source code, and ensuring customer support availability.

Recovery objectives are set pragmatically based on business needs and available resources. We target four-hour recovery time objectives (RTO) for critical development infrastructure and 24-hour RTO for supporting functions. These objectives balance business needs with the cost and complexity of faster recovery capabilities. Recovery point objectives (RPO) are more stringent, with continuous replication for source code and four-hour backup windows for other critical data. These objectives guide our technical architecture and backup strategies.

Continuity procedures are documented in runbooks that provide step-by-step recovery guidance. These procedures assume that key personnel might be unavailable, providing sufficient detail for competent technical staff to execute recovery even without deep system knowledge. We test these procedures quarterly through tabletop exercises and annually through full recovery drills. Testing often reveals gaps or outdated information, driving updates that keep procedures current and effective.



13.2 Security Considerations in Business Continuity

Security cannot be abandoned during crisis response, yet emergency situations often create pressure to bypass normal controls. Our continuity procedures explicitly address security requirements, ensuring that recovery efforts don't create new vulnerabilities. Emergency access procedures provide necessary elevation of privileges while maintaining audit trails and requiring post-event review. We've pre-positioned emergency access credentials in secure storage, accessible to authorized personnel when normal authentication systems are unavailable.

Communication security during incidents receives special attention, as normal channels may be compromised or unavailable. We maintain out-of-band communication methods including personal phone numbers and alternative collaboration platforms. Encryption requirements remain in force for sensitive communications, with pre-shared keys enabling secure communication even if PKI infrastructure is unavailable. These preparations ensure we can coordinate response efforts without exposing sensitive information to attackers who might be monitoring normal channels.

Recovery validation includes security testing to ensure that restored systems haven't been compromised or misconfigured during recovery. We maintain security validation checklists that verify critical controls are properly implemented before systems return to production. This validation has occasionally identified configuration drift that occurred over time, making recovery events valuable for ensuring security baselines are maintained. Post-recovery monitoring watches for signs that incidents might recur or that recovered systems might be targeted.

13.3 Disaster Recovery Architecture

Our disaster recovery architecture leverages cloud-native capabilities to provide resilience without the complexity of traditional disaster recovery sites. Azure's geo-redundant storage automatically replicates critical data to secondary regions, providing protection against regional failures. We've architected our systems for rapid rebuild rather than traditional failover, using infrastructure-as-code to quickly provision new environments when needed.

Source code, our most critical asset, receives multiple layers of protection. Git's distributed nature means every developer has a complete repository copy, providing natural redundancy. Our central repositories in Azure DevOps are geo-replicated, with additional backups to separate Azure storage accounts. Build pipelines can be quickly reconstructed from source, ensuring we can produce deployable software even if build infrastructure is lost. This approach has proven resilient and cost-effective for our needs.



Testing validates our recovery capabilities and identifies gaps before they matter. Monthly backup restoration tests verify that backups are valid and restorable. Quarterly infrastructure rebuild exercises confirm our ability to reconstruct development environments from code. Annual full disaster recovery tests simulate major failures, testing both technical procedures and team coordination. These tests often reveal subtle dependencies or outdated procedures, driving improvements that enhance our real recovery capabilities.

14. Compliance Management

14.1 Legal and Regulatory Compliance

Operating in the software industry requires navigating an evolving landscape of legal and regulatory requirements. Our compliance program takes a risk-based approach, focusing effort where legal obligations are clearest and penalties most severe. We maintain a compliance register that maps applicable laws and regulations to our operations, identifying specific requirements and our methods of compliance. This register is reviewed quarterly and updated as new obligations arise or existing ones change.

Data protection regulations receive particular attention given their proliferation and significant penalties. While our architectural decision to not process customer data significantly reduces our exposure, we still must comply with regulations governing our employee data and business operations. We maintain appropriate privacy notices, implement data subject rights procedures, and ensure cross-border data transfers comply with applicable frameworks. Our privacy-by-design architecture serves as a strong foundation, demonstrating compliance through technical measures rather than just policies.

Working with legal counsel ensures our interpretation of requirements is sound and our compliance measures are adequate. We engage specialized counsel for complex areas like international data transfers or industry-specific regulations. This investment in legal expertise prevents costly misunderstandings and provides confidence in our compliance positions. We document legal consultations and decisions, creating a record that demonstrates due diligence in understanding and meeting our obligations.



14.2 Contractual Security Obligations

Customer contracts often impose security requirements beyond legal minimums, reflecting their own risk management needs. We maintain a library of common contractual security requirements and our standard responses, speeding contract negotiations while ensuring consistency. This library includes evidence of how we meet each requirement, from technical control descriptions to relevant certifications. By proactively addressing common requirements, we reduce negotiation cycles and demonstrate security maturity.

SOC 2 Type II certification serves as objective evidence of our security controls' design and operating effectiveness. This certification covers the Trust Services Criteria relevant to our services: Security, Availability, and Confidentiality. The annual audit process not only provides customer assurance but also drives internal improvements as we address auditor findings. We make our SOC 2 report available under NDA to customers and prospects, providing transparency into our security controls.

Monitoring contractual compliance requires systematic tracking of obligations across all customer agreements. We maintain a compliance matrix that maps specific contractual requirements to responsible parties and evidence of compliance. Regular reviews ensure we remain compliant as contracts evolve through amendments or renewals. When we cannot meet specific requirements, we document risk acceptance or compensating controls, ensuring conscious decisions rather than overlooked obligations.

14.3 Internal Policy Compliance

Policies provide little value if not followed, making compliance monitoring essential. Our monitoring program uses multiple methods to assess compliance, from automated technical scans to management reviews. We avoid creating a police-state atmosphere while ensuring policies are more than just paper exercises. Monitoring focuses on high-risk areas and policies where non-compliance is both likely and impactful.

Automated monitoring leverages our security tools to continuously assess technical compliance. Configuration scanning verifies that systems meet security baselines. Access reviews confirm that permissions align with policy requirements. These automated checks provide broad coverage efficiently, freeing human reviewers to focus on areas requiring judgment. We tune automated monitoring to avoid false positives that would undermine its credibility and usefulness.

Self-assessments by process owners provide insight into operational compliance while building ownership of security requirements. We provide simple questionnaires that



guide assessment without requiring deep security expertise. These assessments often surface practical challenges in following policies, driving improvements that make compliance easier. By involving process owners in compliance assessment, we create partners in security rather than adversaries who view security as imposed from outside.

14.4 Evidence Management

Demonstrating compliance requires evidence that controls are not just designed but operating effectively. Our evidence management approach balances thoroughness with practicality, focusing on evidence that provides meaningful assurance without creating overwhelming administrative burden. We leverage automated evidence collection where possible, such as system-generated logs and reports that demonstrate control operation without manual effort.

Evidence retention periods reflect both compliance requirements and practical limitations. We maintain audit logs for one year, providing sufficient history for most compliance needs while avoiding excessive storage costs. Other evidence types have retention periods based on specific requirements or business needs. Clear retention schedules ensure consistent treatment and prevent both premature deletion and indefinite retention of evidence that no longer provides value.

Audit support procedures ensure smooth interactions with auditors while protecting business operations. We maintain standard evidence packages for common audit requests, reducing preparation time and ensuring consistency. Designated audit liaisons coordinate requests and responses, preventing auditors from disrupting operations through uncoordinated requests to multiple staff members. Post-audit reviews capture lessons learned and identify evidence gaps for proactive remediation before the next audit cycle.

Policy development follows a collaborative process that balances security requirements with operational practicality. Draft policies undergo review by affected stakeholders who provide feedback on feasibility and potential unintended consequences. This collaborative approach improves policy quality while building buy-in from those who must implement the policies. We've learned that policies developed in isolation by security staff often face resistance or prove impractical, while collaborative development produces policies that are both secure and workable.

Version control and change tracking ensure policy modifications are transparent and auditable. We maintain complete version history showing what changed, when, and why. This history proves valuable during audits and helps new employees understand policy evolution. Major policy changes require executive approval, ensuring leadership



awareness and support. Minor clarifications follow a simplified process to avoid bureaucratic delays while maintaining appropriate oversight.

16.2 Policy Communication and Accessibility

The best policies are worthless if employees can't find or understand them. Our policy repository provides centralized access to all current policies, procedures, and standards. Simple navigation and robust search capabilities help employees quickly find relevant information. We avoid burying policies in complex document management systems, instead using familiar tools that employees access daily. Mobile-friendly formatting ensures policies are accessible to our remote workforce regardless of device.

Plain language communication makes policies understandable to all employees, not just security professionals. We write policies in clear, concise language, avoiding jargon where possible and defining technical terms when necessary. Examples and scenarios illustrate abstract concepts, making requirements concrete and relatable. Policy summaries provide quick reference for key points, with links to full policies for those needing more detail. This layered approach serves both quick reference needs and comprehensive understanding.

Change communication ensures employees are aware of new or updated policies that affect them. We use multiple channels including email announcements, team meeting discussions, and internal wiki updates. Major changes receive dedicated training sessions to ensure understanding. We track acknowledgment of critical policy changes, ensuring all affected employees are aware of new requirements. This proactive communication prevents the excuse of ignorance and ensures consistent implementation across the organization.

16.3 Exception Management

No policy can anticipate every situation, making a robust exception process essential for handling edge cases without compromising security. Our exception management framework provides clear pathways for requesting, evaluating, and tracking policy exceptions while maintaining security integrity and audit trails.

16.3.1 Exception Request Process

All policy exception requests must follow our standardized process to ensure consistent evaluation and documentation:

Request Submission Requirements:



1. **Business Justification:** Clear explanation of why the exception is needed
2. **Specific Policy Reference:** Exact policy section(s) requiring exception
3. **Duration:** Requested timeframe (maximum 12 months)
4. **Risk Assessment:** Potential security impacts of granting the exception
5. **Compensating Controls:** Alternative security measures to mitigate risks
6. **Affected Systems/Data:** Scope of systems or data classifications involved

Submission Process:

- Submit via IT ticketing system with "Policy Exception Request" category
- Include all required information using standard exception template
- Attach supporting documentation (project plans, vendor requirements, etc.)
- Request must be submitted by employee's manager or above

16.3.2 Risk Assessment and Approval Authority

Our risk-based approval matrix ensures appropriate oversight while avoiding unnecessary bureaucracy:

Risk Classification Criteria:

- **Low Risk:** Limited scope, no customer data involved, strong compensating controls
- **Medium Risk:** Broader scope, internal sensitive data, temporary business need
- **High Risk:** Customer-facing systems, regulatory implications, extended duration
- **Critical Risk:** Core security principles compromised, no viable compensating controls

Approval Authority Matrix:

Risk Level	Approval Required	Maximum Duration	Review Frequency
Low	Department Manager	6 months	Quarterly
Medium	IT Manager	6 months	Monthly
High	CTO	3 months	Monthly
Critical	CEO	1 month	Weekly

Evaluation Timeline:

- Low/Medium Risk: Decision within 3 business days
- High/Critical Risk: Decision within 5 business days



- Emergency exceptions: Verbal approval within 4 hours, documentation within 24 hours

16.3.3 Exception Documentation and Tracking

All approved exceptions are formally documented and actively tracked:

Exception Record Requirements:

- Unique exception ID (EXC-YYYY-###)
- Approval date and approvers
- Expiration date (with calendar reminders at 30, 15, and 5 days)
- Risk rating and accepted residual risks
- Compensating controls implemented
- Quarterly review notes
- Link to original request and supporting documents

Exception Register Management:

- Centralized register maintained in IT service management platform
- Monthly summary reports to leadership team
- Quarterly detailed review by CTO
- Annual audit by external assessor
- Automated notifications for expiration and review dates

16.3.4 Emergency Exception Process

Critical business needs sometimes require immediate policy exceptions:

1. **Immediate Verbal Approval:** Contact CTO or CEO for verbal authorization
2. **Temporary Controls:** Implement agreed compensating controls immediately
3. **24-Hour Documentation:** Submit formal exception request within one business day
4. **72-Hour Risk Assessment:** Complete full risk evaluation within three days
5. **30-Day Resolution:** Convert to standard exception or remediate within 30 days

Emergency exceptions bypassing this process will be treated as security incidents.

16.3.5 Exception Monitoring and Metrics

We track exception patterns to identify systemic issues and policy improvement opportunities:

Key Metrics:



- Exception volume by policy area
- Average exception duration
- Repeat exception requests
- Expired exception cleanup rate
- Emergency exception frequency

Quarterly Analysis:

- Identify policies with high exception rates
- Evaluate if policy updates would reduce exceptions
- Review effectiveness of compensating controls
- Assess trends in risk acceptance

Continuous Improvement:

When multiple similar exceptions occur, we evaluate whether the underlying policy needs revision rather than continuing to grant exceptions. This prevents policy drift and ensures our security framework remains practical and effective.

16.4 Governance and Oversight

Effective governance ensures our security program remains aligned with business objectives while maintaining appropriate independence. Our governance structure reflects our company size, avoiding heavyweight processes while ensuring appropriate oversight. The CTO reports directly to the CEO on security matters, providing independence from operational pressures. Regular board updates ensure director-level awareness of security posture and significant risks.

Security metrics provide objective assessment of program effectiveness. We track meaningful indicators like incident frequency and impact, vulnerability remediation timelines, and training completion rates. These metrics are presented in business context, showing how security supports operational goals rather than presenting abstract technical measurements. Trend analysis identifies improving or degrading security posture, enabling proactive intervention before minor issues become major problems.

External validation through audits and assessments provides independent perspective on our security program. Annual SOC 2 audits offer third-party validation of control effectiveness. Periodic penetration tests challenge our technical controls. These external assessments often identify blind spots that internal reviews miss. We treat findings as improvement opportunities rather than failures, using them to strengthen our security program. This commitment to continuous improvement based on



independent assessment demonstrates security program maturity despite our small size.

17. Enforcement and Accountability

17.1 Creating a Culture of Security Accountability

Enforcement of security policies requires a delicate balance between ensuring compliance and maintaining a positive work environment. Our approach emphasizes education and support over punishment, recognizing that most policy violations result from misunderstanding or oversight rather than malicious intent. By creating an environment where employees feel safe asking questions and reporting mistakes, we achieve better security outcomes than through fear-based enforcement.

Accountability begins with clear expectations communicated during onboarding and reinforced through regular training. Every employee signs acknowledgment of security policies, confirming understanding of their responsibilities. More importantly, we ensure employees have the knowledge and tools needed to meet these expectations. When violations occur, our first response is to understand why—was it lack of knowledge, inadequate tools, or conflicting priorities? Addressing root causes prevents recurrence more effectively than punishment alone.

Progressive discipline provides a measured response to violations while offering opportunities for improvement. First violations typically result in additional training and coaching, ensuring the employee understands both the requirement and its importance. Repeated violations escalate through formal warnings to potential termination, but we rarely reach severe consequences because our supportive approach resolves most issues early. This progressive approach documents patterns for truly problematic employees while helping well-intentioned employees succeed.

17.2 Monitoring and Detection

Effective enforcement requires knowing when violations occur, but monitoring must respect employee privacy and maintain trust. Our monitoring focuses on security-relevant activities rather than general employee surveillance. We're transparent about what we monitor and why, helping employees understand that monitoring protects both company assets and their own reputation. Technical controls automate much monitoring, providing consistent enforcement without human bias.

Behavioral monitoring identifies patterns that might indicate security issues or policy violations. Unusual access patterns, like accessing systems outside normal working hours or downloading large data volumes, trigger alerts for investigation. We tune



these alerts to minimize false positives while catching genuinely concerning activities. Investigation procedures ensure fair, consistent evaluation of alerts, distinguishing legitimate unusual activities from actual security concerns.

Regular compliance assessments complement continuous monitoring by examining areas that automated tools might miss. These assessments might include physical security walk-throughs in home offices (with appropriate notice and consent), review of security configurations on devices, or verification of training completion. We position these assessments as health checks that help employees maintain security rather than gotcha exercises looking for violations. This supportive approach encourages cooperation and honest discussion of challenges.

17.3 Consistent and Fair Response

When violations are confirmed, our response must be consistent, fair, and proportional to the violation's severity and impact. We maintain a violation response matrix that guides decisions based on violation type, intent, impact, and history. This matrix ensures similar violations receive similar responses regardless of who is involved. Documentation of all violations and responses creates a record that demonstrates consistent treatment and supports any necessary formal disciplinary action.

Mitigating factors receive appropriate consideration in our response. Self-reporting of violations demonstrates integrity and reduces response severity. Immediate action to minimize impact shows good judgment even after a mistake. Cooperation with investigation and remediation efforts indicates learning from the incident. These factors don't eliminate consequences but can reduce them, encouraging positive behaviors we want to reinforce. Conversely, attempts to hide violations or refusal to cooperate escalate responses.

Learning from violations improves our overall security posture. Post-violation reviews examine not just individual fault but systemic factors that enabled the violation. Did policies clearly communicate requirements? Were appropriate tools available? Were there conflicting priorities that made compliance difficult? These reviews often identify improvement opportunities for policies, training, or tools. By treating violations as learning opportunities for the organization, not just the individual, we strengthen our security program while maintaining morale.



18. Related Policies and Standards

18.1 Integrated Policy Framework

The Information Security Policy does not stand alone but serves as the cornerstone of an integrated framework of policies, standards, and procedures that collectively protect OversightAI's assets. Each supporting document addresses specific aspects of security in greater detail while maintaining alignment with the principles and requirements established here. This hierarchical approach provides clarity and prevents conflicts while allowing detailed guidance where needed.

Our Access Control Policy expands on the access management principles outlined in Section 5, providing detailed requirements for authentication, authorization, and account management. It specifies technical standards for password complexity, multi-factor authentication implementation, and privileged access management. Procedures for account provisioning, modification, and termination ensure consistent implementation across all systems. This detailed guidance enables IT staff to implement controls correctly while maintaining flexibility for system-specific requirements.

The Data Classification and Handling Policy elaborates on the classification framework introduced in Section 4, providing specific handling requirements for each classification level. It addresses data throughout its lifecycle from creation through disposal, ensuring appropriate protection at each stage. Detailed marking, transmission, storage, and disposal requirements prevent confusion about how to handle different information types. This policy is particularly important for our distributed workforce, providing clear guidance for protecting information in home office environments.

18.2 Operational Standards and Procedures

Supporting standards translate policy requirements into specific technical implementations. Our Secure Development Standards provide language-specific guidance for avoiding common vulnerabilities. Rather than abstract security principles, these standards show actual code examples of vulnerable and secure implementations. Developers can reference these standards during coding, reducing security defects through proactive guidance rather than reactive finding and fixing.

Incident Response Procedures operationalize the framework described in Section 12, providing step-by-step instructions for detecting, responding to, and recovering from security incidents. These procedures include contact lists, escalation criteria, and evidence preservation techniques. Checklists ensure critical steps aren't missed during the stress of incident response. Regular drills validate these procedures and identify



necessary updates as our environment evolves.

Change Management Procedures ensure the principles in Section 10 are consistently applied to all system modifications. From emergency patches to major architectural changes, these procedures scale appropriately while maintaining necessary controls. Template forms capture required information consistently while review checklists ensure security considerations aren't overlooked. Automation enforces many procedural requirements, reducing human error while improving efficiency.

18.3 Compliance and Audit Documentation

The Risk Management Policy provides the framework for identifying, assessing, and treating the risks that our security controls address. It establishes risk appetite, assessment methodologies, and treatment strategies that guide security investment decisions. Regular risk assessments identify emerging threats and evaluate control effectiveness, ensuring our security program evolves appropriately. This risk-based approach ensures we focus limited resources where they provide the greatest benefit.

Business Continuity and Disaster Recovery Plans ensure security remains effective even during disruptions. These plans address maintaining security controls during incidents, recovering securely from disasters, and validating security after recovery. They recognize that emergencies create pressure to bypass security controls and provide pre-approved procedures that maintain necessary protection while enabling rapid response. Regular testing validates these plans and identifies necessary updates.

The Human Resources Security Policy addresses security throughout the employee lifecycle from recruitment through termination. It establishes background check requirements, security training obligations, and termination procedures that protect company assets. This policy is particularly important for our remote workforce, addressing unique challenges like equipment recovery and access revocation when employees may be geographically distant. Clear procedures ensure consistent handling regardless of circumstances.

18.4 Continuous Improvement Through Integration

Policy integration enables continuous improvement across our security program. Lessons learned from incidents inform updates to multiple policies, ensuring comprehensive remediation. New threats identified through risk assessment drive updates to technical standards and training programs. Audit findings might reveal gaps requiring new procedures or policy clarifications. This integrated approach ensures improvements strengthen our entire security program rather than creating point solutions.



Regular cross-policy reviews identify conflicts, gaps, and redundancies that naturally develop as policies evolve independently. These reviews ensure policies remain mutually reinforcing rather than contradictory. They also identify opportunities to simplify by consolidating related requirements or eliminating outdated provisions. This maintenance prevents policy bloat that would make compliance difficult and reduce policy credibility.

Stakeholder feedback on policy interactions helps identify practical challenges in implementation. When employees report difficulty understanding how multiple policies apply to specific situations, we create integrated guidance that clarifies requirements. These real-world scenarios often reveal policy conflicts or gaps that weren't apparent during individual policy development. By maintaining an integrated view of our policy framework, we ensure it remains practical and effective for protecting OversightAI's assets while enabling business success.

19. Definitions and Glossary

Access Control: The process of granting or denying specific requests to obtain and use information and related information processing services.

Asset: Any item of value to the organization, including information, systems, physical items, and intangible items like reputation.

Authentication: The process of verifying the identity of a user, device, or system, often as a prerequisite to allowing access to resources.

Authorization: The process of determining what permissions an authenticated user has within a system or application.

Availability: Ensuring timely and reliable access to information by authorized users when needed.

Business Continuity: The capability of the organization to continue delivery of products or services at acceptable predefined levels following a disruptive incident.

Confidentiality: Ensuring that information is not made available or disclosed to unauthorized individuals, entities, or processes.

Data Classification: The process of organizing data into categories based on sensitivity level and required protection.



Encryption: The process of converting information into a code to prevent unauthorized access, ensuring confidentiality and integrity.

Incident: Any event that has the potential to affect the confidentiality, integrity, or availability of information or information systems.

Information Security: The practice of protecting information by mitigating information risks through implementing appropriate controls.

Integrity: The accuracy and completeness of information and processing methods, ensuring information has not been modified or destroyed in an unauthorized manner.

Least Privilege: The principle of providing users only the minimum levels of access needed to perform their job functions.

Multi-Factor Authentication (MFA): Authentication requiring two or more verification factors (something you know, have, or are).

Risk: The potential for loss, damage, or destruction of an asset as a result of a threat exploiting a vulnerability.

Risk Assessment: The process of identifying, analyzing, and evaluating risks to determine appropriate treatment strategies.

Security Control: A safeguard or countermeasure designed to avoid, detect, counteract, or minimize security risks.

Security Event: Any observable occurrence in a system or network that may indicate a security incident.

Threat: Any circumstance or event with the potential to adversely impact organizational operations, assets, or individuals.

Vulnerability: A weakness in a system, application, or process that could be exploited by a threat to gain unauthorized access or cause harm.

Zero Trust: A security model that requires strict verification for every person and device trying to access resources, regardless of location.

20. Document Control

NIST Controls: PM-4, SA-5



Version	Date	Author	Changes
1.0	January 1, 2025	CTO	Initial comprehensive version
2.0	Jun 25, 2025	CTO	Added NIST control mappings throughout document and new Appendix G

Review and Approval

- Prepared By: _____ Date: _____
- Approved By: _____ Date: _____

Next Review Date: January 1, 2026

Distribution:

- All Employees: Via company policy portal
- IT Team: Direct distribution for implementation
- Executive Team
- External Auditors (upon request)

21. Appendices

Appendix A: NIST Control Mapping

This appendix provides a comprehensive mapping of Information Security Policy sections to SOC 2 Trust Services Criteria and NIST Cybersecurity Framework controls. This mapping assists auditors in identifying control coverage and helps OversightAI ensure comprehensive compliance.



Control Mapping Matrix

Policy Section	SOC 2 Criteria	NIST CSF	Control Objective
1. Purpose and Scope	CC1.1, CC1.2	ID.GV-1	Organizational security governance and scope
2. Information Security Principles	CC1.4, CC2.1	ID.GV-1, ID.GV-3	Security philosophy and commitment
3. Roles and Responsibilities	CC1.3, CC1.5	ID.GV-2, ID.AM-6	Clear accountability and ownership
4. Asset Management	CC3.4, CC6.4	ID.AM-1, ID.AM-2	Inventory and classification of assets
5. Access Control	CC6.1, CC6.2, CC6.3	PR.AC-1, PR.AC-4	Identity management and authorization
5.2 Identity and Authentication	CC6.1	PR.AC-1, PR.AC-7	Strong authentication requirements
5.3 Authorization and Privileged Access	CC6.3	PR.AC-4	Least privilege and segregation
5.4 Remote Access Security	CC6.6	PR.AC-3, PR.AC-5	Secure remote connectivity
6. Cryptography and Data Protection	CC6.1, CC6.7	PR.DS-1, PR.DS-2	Encryption and key management
7. Physical and Environmental Security	CC6.4, CC6.5	PR.AC-2, PR.DS-3	Physical access and environmental controls
8. Operations Security	CC7.1, CC7.2	DE.AE-1, DE.CM-1	Operational control and monitoring
8.1 Customer Data Isolation	CC6.1, P3.2	PR.DS-5	Architectural security controls
8.2 Change Management	CC8.1	PR.IP-3	Change control processes



8.3 Security Monitoring	CC7.1	DE.AE-1, DE.CM-1	Event detection and monitoring
8.4 Vulnerability Management	CC7.1	ID.RA-1, RS.MI-3	Vulnerability identification and remediation
9. Communications Security	CC6.6, CC6.7	PR.AC-5, PR.DS-2	Network and application security
10. System Development	CC8.1	PR.IP-2, PR.IP-3	Secure development lifecycle
11. Supplier and Third-Party Security	CC9.1, CC9.2	ID.SC-1, ID.SC-2	Vendor risk management
12. Incident Management	CC7.3, CC7.4, CC7.5	RS.RP-1, RS.CO-2	Incident response and recovery
13. Business Continuity	A1.1, A1.2, A1.3	PR.IP-9, RC.RP-1	Resilience and recovery planning
14. Compliance Management	CC2.2, CC2.3	ID.GV-3, ID.GV-4	Legal and regulatory compliance
15. Security Training	CC1.4	PR.AT-1, PR.AT-2	Security awareness and training
16. Policy Management	CC5.2	ID.GV-1	Policy lifecycle and governance
16.3 Exception Management	CC3.3, CC5.3	ID.RA-3	Risk acceptance and exception handling
17. Enforcement and Accountability	CC1.5	PR.IP-11	Disciplinary process and accountability

SOC 2 Common Criteria (CC) Coverage Summary

Control Environment (CC1)

- CC1.1-CC1.5: Fully addressed through governance structure, roles, responsibilities, and accountability measures

Communication and Information (CC2)

- CC2.1-CC2.3: Covered by security principles, compliance management, and policy communication



Risk Assessment (CC3)

- CC3.1-CC3.4: Addressed via risk management integration, asset management, and exception processes

Monitoring Activities (CC4)

- CC4.1-CC4.2: Implemented through continuous monitoring and security metrics

Control Activities (CC5)

- CC5.1-CC5.3: Established through policy framework, change control, and exception management

Logical and Physical Access (CC6)

- CC6.1-CC6.8: Comprehensive coverage through access control, encryption, and security architecture

System Operations (CC7)

- CC7.1-CC7.5: Covered by monitoring, vulnerability management, and incident response

Change Management (CC8)

- CC8.1: Addressed in change management and SDLC security

Risk Mitigation (CC9)

- CC9.1-CC9.2: Implemented through vendor management and third-party security

NIST Cybersecurity Framework Alignment

Identify (ID)

- Asset Management, Risk Assessment, Governance, and Supply Chain covered

Protect (PR)

- Access Control, Awareness Training, Data Security, and Protective Technology implemented

Detect (DE)

- Anomalies and Events, Continuous Monitoring addressed

Respond (RS)

- Response Planning, Communications, and Mitigation covered

Recover (RC)

- Recovery Planning and Communications through Business Continuity



Using This Mapping

For Auditors:

1. Use this matrix to quickly identify where specific controls are documented
2. Cross-reference with evidence requests to locate relevant procedures
3. Verify control implementation through referenced policy sections

For OversightAI Staff:

1. Reference during control implementation to ensure compliance
2. Use for gap analysis when requirements change
3. Guide for evidence collection during audit preparation

Maintenance:

- Review quarterly during policy check-ins
- Update when policy sections change
- Validate against SOC 2 criteria updates annually
- Enhance based on auditor feedback

Acknowledgment

By signing below, I acknowledge that I have read, understood, and agree to comply with the OversightAI Information Security Policy and all related security policies and procedures. I understand that violation of these policies may result in disciplinary action up to and including termination of employment or contract.

I further acknowledge that:

- I will protect OversightAI's confidential information and assets
- I will report any suspected security incidents or policy violations immediately
- I will complete all required security training
- I will maintain the security of any devices and accounts assigned to me
- I understand that my activities on company systems may be monitored for security purposes

Employee Name: _____

Employee Signature: _____

Date: _____

Manager Name: _____



Manager Signature: _____

Date: _____

This acknowledgment must be completed by all employees upon hire and annually thereafter, with signed copies maintained in personnel files.